
ARM Architecture



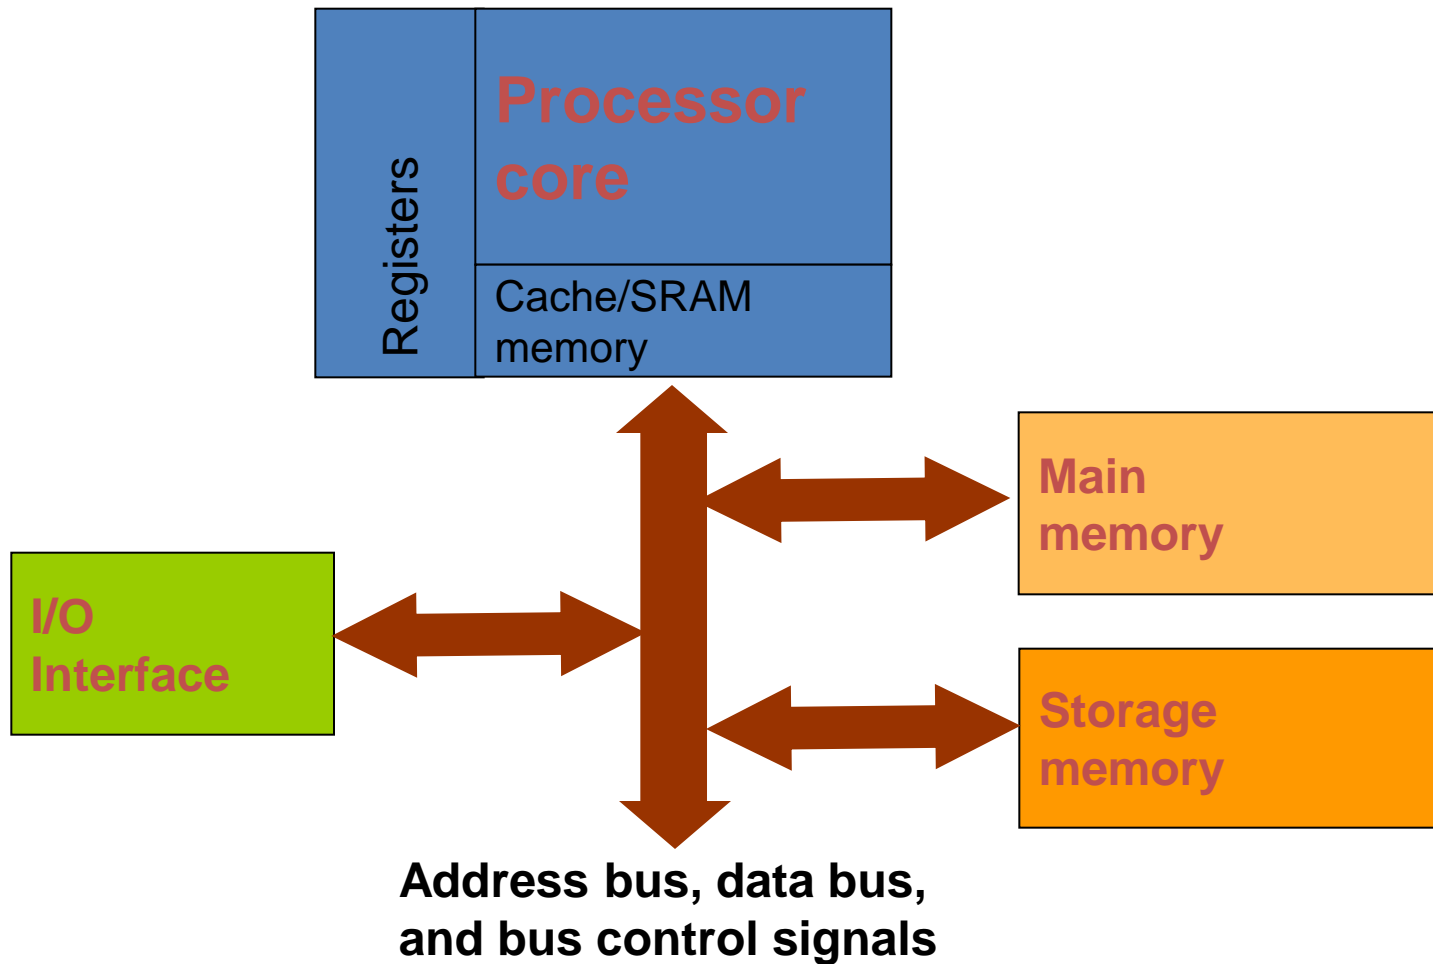
The RISC Philosophy

- Original RISC design (e.g. MIPS)
 - aims for high performance through
 - reduced number of instruction classes
 - large general-purpose register set
 - load-store architecture
 - fixed length instructions
 - pipelines
 - enables simpler hardware, therefore enabling it to scale to higher operating frequencies

A Processor Architecture Refresher



Basic Processor-Based System



System Components

- The basic components:
 - Processor with its associated temporary memory (registers and cache if available) for code execution
 - Main memory and secondary memory where code and data are temporarily and permanently stored
 - Input and output modules that provide interfaces between the processor and the user
- Connected through an interface bus that consists of Address, Data, and Control signals
 - e.g., AMBA bus for the ARM-based processor

Memory Hierarchy


- A typical processor is supported by:
 - on-board main memory (e.g. SDRAM up to GB)
 - on-chip or on-die cache memory (e.g. SRAM KB to MB)
 - on-die registers
- Some processors also provide general purpose on-chip
 - SRAM (e.g. embedded processor) which may be configured as SRAM/Cache combination (e.g. TI's DSP)
- Typically, a processor also utilizes secondary non-volatile memory
 - for permanent code and data storage like Flash-based memory and hard disks

Address Space

- The address space of a processor depends on its address decoding mechanism
 - Its size will depend on the number of address bits used
- Depending on the processor design, there may be two types of address space
 - one is used by normal memory access
 - another one is reserved for I/O peripheral registers (control, status, and data)
 - need extra control signal or special means of accessing the alternate address space

Address Space (con't)

- Refer to the range of addresses that can be accessed by the processor, determined by the number of address bit utilized in the processor architecture
- Some processor families (e.g. ARM) utilize only one address space for both memory and I/O devices
 - i.e. everything is mapped in the same address space

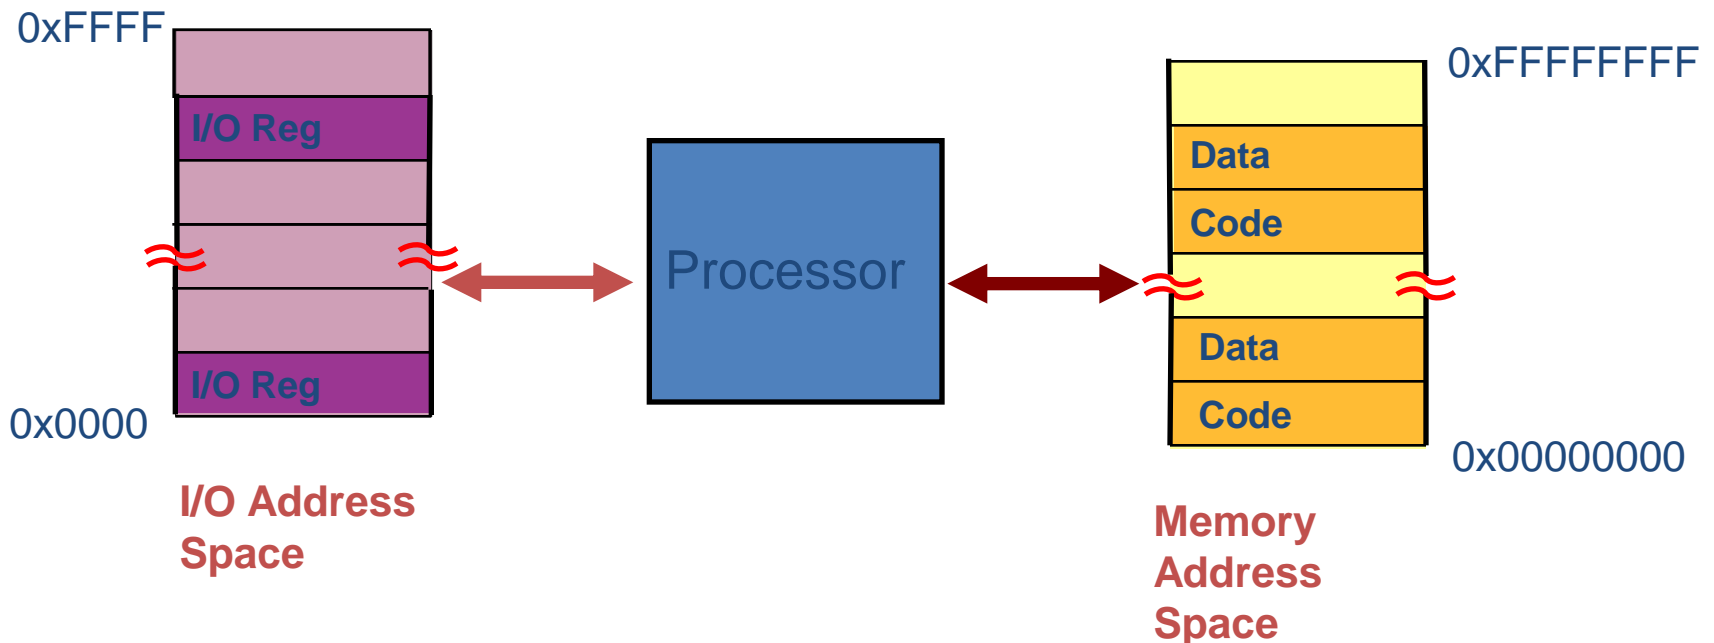


Memory Mapped vs I/O Mapped

Some processor families have two address spaces.

E.g., for the x86 processor, memory and I/O devices can be mapped in two different address spaces:

- memory address space and I/O address space

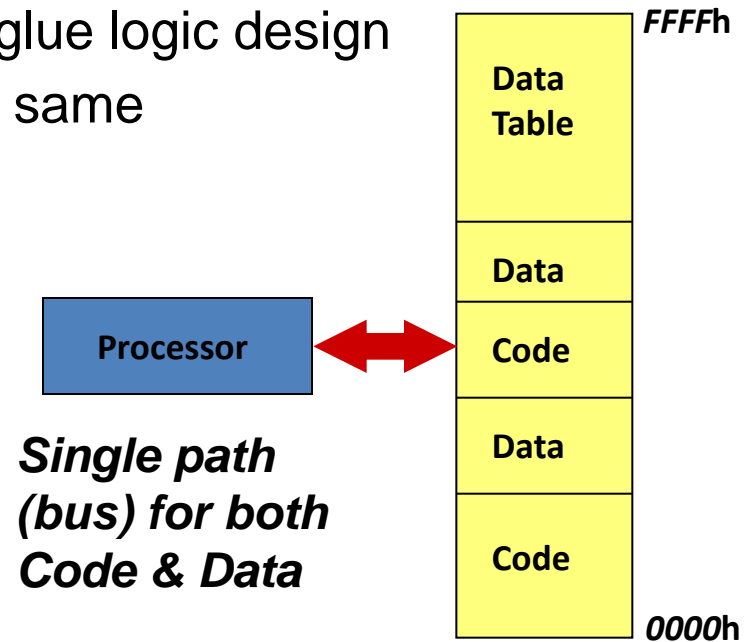


Memory System Architectures

- Two types of information are found in a typical program code:
 - i. Instruction codes for execution
 - ii. Data that is used by the instruction codes
- Two classes of memory systems designed to store the information:
 - i. von Neumann architecture
 - ii. Harvard architecture

von Neumann Architecture

- The von Neumann architecture utilizes only one memory bus for both instruction fetching and data access
 - simplifies the hardware and glue logic design
 - code and data located in the same address space

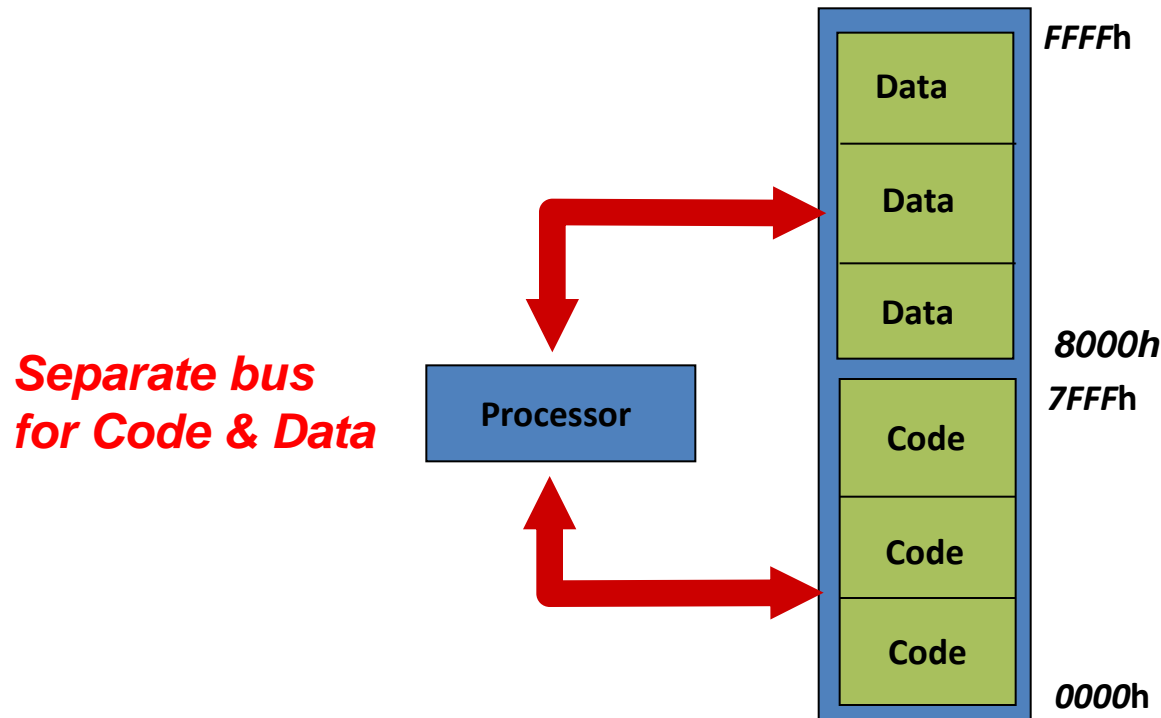


von Neumann Features

- Single memory interface bus
 - simplifies the hardware and glue logic design
- More efficient use of memory
 - code and data can reside in the same physical memory chip
- More flexible programming style
 - e.g., can include self-modified code
- But data may overwrite code (e.g. due to program bug)
 - need memory protection (e.g. hardware-based MPU)
- Bottleneck in code and data transfer
 - only one memory bus for both data and code fetching

Harvard Architecture

- The Harvard architecture utilizes separate instruction bus and data bus
 - code and data may still share the same memory space




Harvard Features

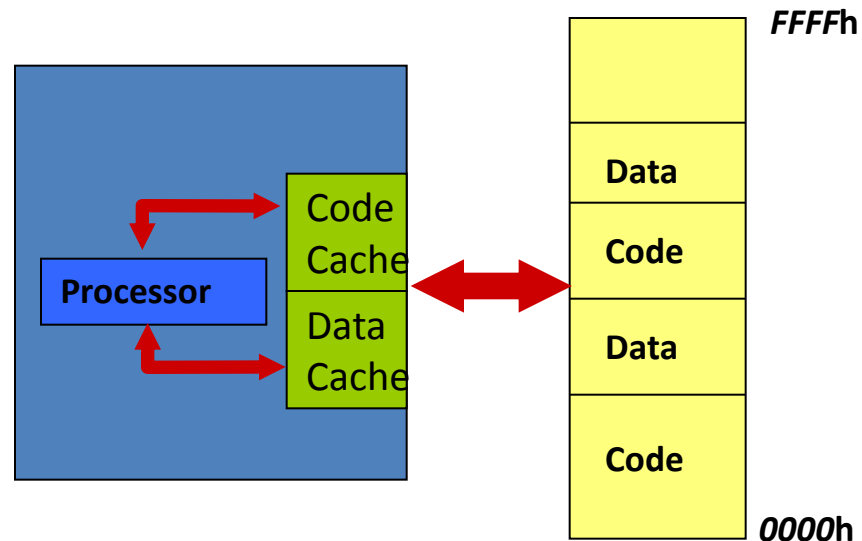
- Separate instruction and data buses
 - allow code and data access at the same time which gives improved performance
 - provide better support for instruction pipeline operations and shorter instruction execution time
 - allow different sizes of data and instructions to be used which results in more flexibility
 - do not incur any code corruption by data which makes the operations more robust
- But more sophisticated hardware glue logic is required to support multiple interface buses
- Cortex-M3 core is based on the Harvard architecture with separate buses for instructions and data

Architecture Variations

Independent data and code memory but with one shared bus (e.g. 8051)



Two separate internal bus for code & data (e.g. ARM9)

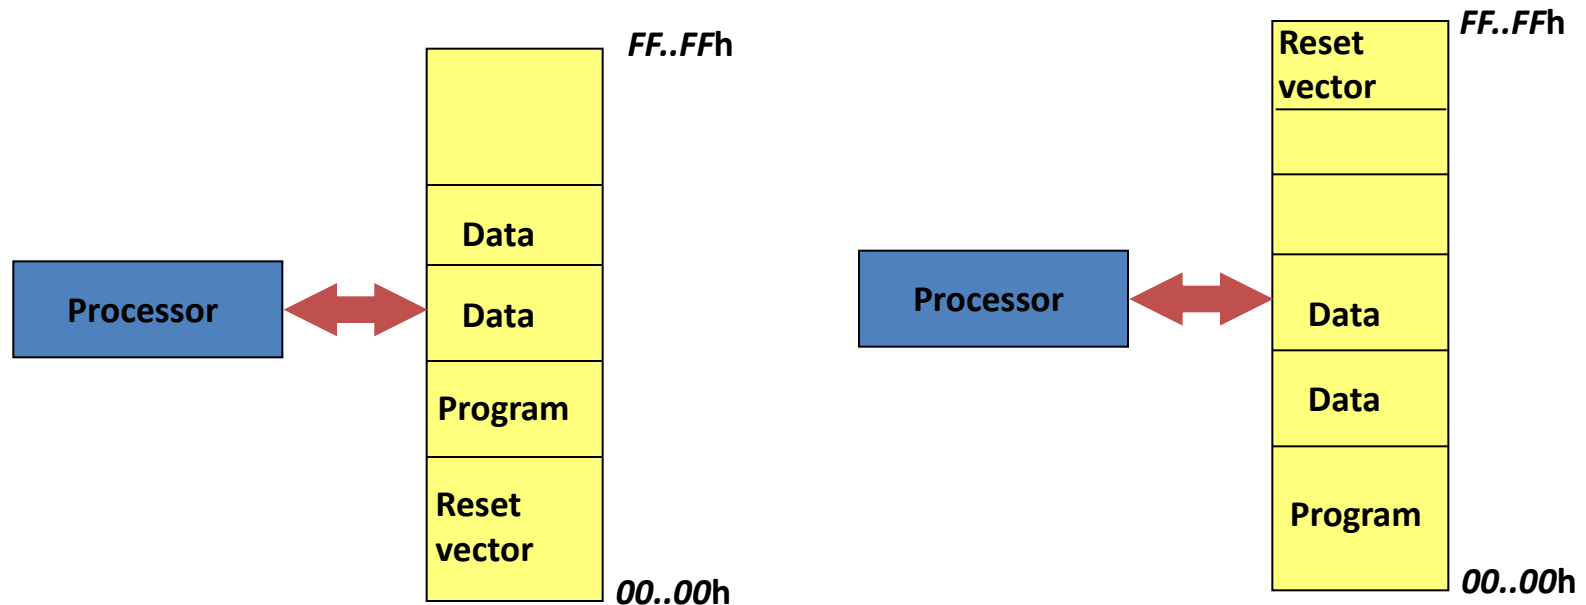


Top Boot and Bottom Boot

Different processor families use different locations for their reset vectors for boot-up.

Examples:

- x86 processors boot up from the top of the memory space
- ARM processors boot up from the bottom of its memory space



Processor 'Size'

- Processor size is described in terms of 'bits' (e.g. an 8-bit or 32-bit processor)
 - corresponds to the data size that can be manipulated at a time by the processor
 - typically reflected in the size of the processor (internal) data path and register bank
- An 8-bit processor can only manipulate one byte of data at a time, while a 32-bit processor can handle one 32-bit double word sized data at a time
 - even though the data content may only be of single byte size

Registers

- The most fundamental storage area in the processor
 - is closely located to the processor
 - provides very fast access, operating at the same frequency as the processor clock
 - but is of limited quantity (typically less than 100)
- Most are of the general purpose type and can store any type of information:
 - data – e.g., timer value, constants
 - address – e.g., ASCII table, stack
- Some are reserved for specific purposes
 - program counter (r15 in ARM)
 - program status register (CPSR in ARM)


Data Organization in Memory

- Typically, memory contains a storage locations that can store data of a certain fixed size
 - most commonly of the 8-bit (byte) size
- Each location is provided with a unique address.
- Depending on the data path size of the processor
 - the memory content is accessible in the size of an 8-bit byte, a 16-bit half word, a 32-bit word, and even a 64-bit double word

Data Alignment

- A 32-bit data consists of four bytes of data, which are stored in four successive memory locations
- Data and code must be aligned to the respective address size boundary.
 - e.g., for a 32-bit system which aligns to the word boundary, the lowest two address bits equal to zero
- But what is the order of the four bytes of data?
 - depends on the Endianness adopted
- In the Little Endian format,
 - the least significant byte (LSB) is stored in the lowest address of the memory, with the most significant byte (MSB) stored in the highest address location of the memory.
- In the Big Endian format,
 - the least significant byte (LSB) is stored in the highest address of the memory, with the most significant byte (MSB) stored in the lowest address location of the memory.

Data Endianness



Comparison

- Little Endian
 - The order matched with processor instructions typically process numbers from LSB to MSB
 - The byte number corresponds with the address offset, suitable for multi-precision data manipulation
- Big Endian
 - Can compare numerical data by just accessing the zero offset byte
 - Corresponds to the written order of number (starting with the most significant digit)
- Some processors (e.g. ARM) have bi-endian hardware that feature 'switchable' endianness

CISC

- Features of the Complex Instruction Set Computing (CISC):
 - many instructions
 - complex instructions
 - each instruction can execute several low level operations
 - complex addressing modes
 - smaller number of registers needed
- A semantically rich instruction set is accommodated by allowing instructions of variable length

Advantages of CISC

- As each instruction can execute several low level operations,
 - the code size is reduced to save on memory requirements
 - less main memory access is required and hence processing time is reduced (faster)
- Backward code compatibility is maintained
 - can add new (and more powerful) instructions while retaining the 'old' instruction set for code compatibility (i.e. legacy programs can still run)
- Easy to program
 - direct support of high-level language constructs
 - complex instructions that fit well with high-level language expressions

Limitations of CISC

- A highly encoded instruction set needs to be decoded by hardwired microcode electronic circuitry
 - more complex hardware design
 - slower instruction decoding/execution
- Variable length instructions
 - different execution time among instructions
 - affects pipelined operations

RISC

RISC – Reduced Instruction Set Computing

- Small instruction sets
- Simpler instructions
- Fixed length instructions
- Large number of registers
- Simpler addressing mode with the Load/Store instructions for accessing memory

Advantages of RISC

- Simpler instructions
 - one clock per instruction gives faster execution than on a CISC processor with the same clock speed
- Simpler addressing mode
 - faster decoding
- Fixed length instructions
 - faster decoding and better pipeline performance
- Simpler hardware
 - less silicon area
 - less power consumption

RISC Memory Footprint

- The RISC processor typically needs more memory than a CISC processor does to store the same program
 - complex functions performed in a single but slower instruction in a CISC processor may require two, three, or more simpler instructions in a RISC processor
- To reduce memory requirements and cost
 - ARM provides the 16-bit Thumb instruction set as an option for its RISC processor cores

Limitations of RISC

- Fewer instructions than CISC
 - as compared to CISC, RISC needs more instructions to execute one task
 - code density is less
 - needs more memory
- No complex instructions
 - no hardware support for division or floating-point arithmetic operations
 - needs a more complex compiler and longer compiling time
- But ARM also adds DSP-like instructions to support commonly used signal processing functions

Instruction Code Format

- Opcode encoding depends on the number of bits used
 - Example: For ARM, all instructions are of 32-bit length, but only 8 bits (bit 20 to 28) are used to encode the instruction. Hence a total of $2^8 = 256$ different instructions are possible
- A typical instruction is encoded with a specific bit pattern that consists of the following:
 - an opcode field specifying the operation to be performed
 - an operand(s) identification (address) field that depends on the modes of addressing;
 - this provides the address of the register/memory location (s) that store the operand(s), or the operand itself.

Instruction Opcode Types

General categories of instruction operations:

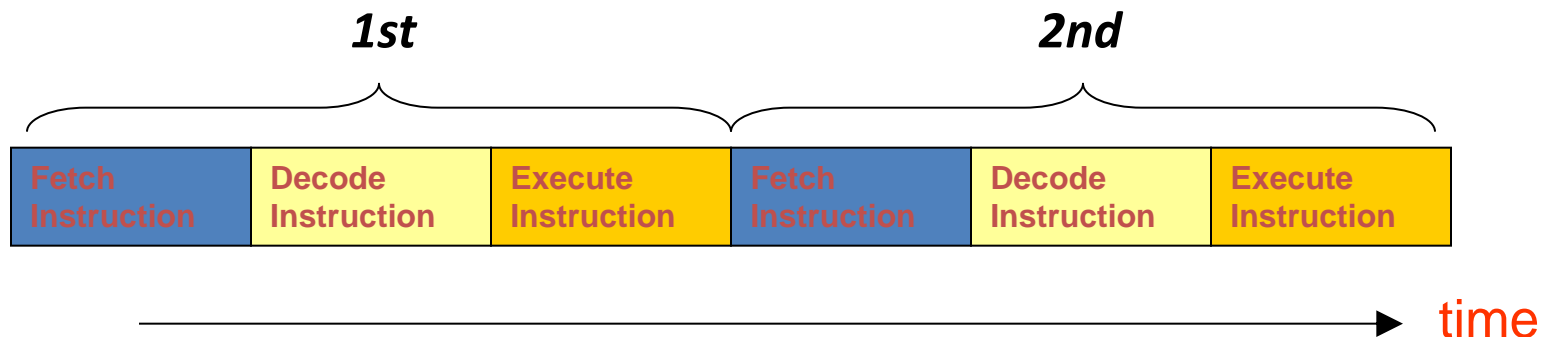
- Data transfer
e.g., move, load, and store
- Data manipulation
e.g., add, subtract, logical operation
- Program control
e.g., branch, subroutine call

Operand Addressing Types

- Immediate addressing
 - operand is given in the instruction
- Register addressing
 - operand is stored in a register
- Direct addressing
 - operand is stored in memory, with the address given in the instruction
- Indirect (Index) addressing
 - operand is stored in memory, with the address given in a register (address and an offset given in the instruction)
- Implied addressing
 - implicit location like stack and program counter

Instruction Execution

- Multiple stages are involved in executing an instruction.
 - Example:
 - 1) Fetching the instruction code
 - 2) Decoding the instruction code
 - 3) Executing the instruction code
- Hence multiple processor clock cycles are needed to execute one single instruction.

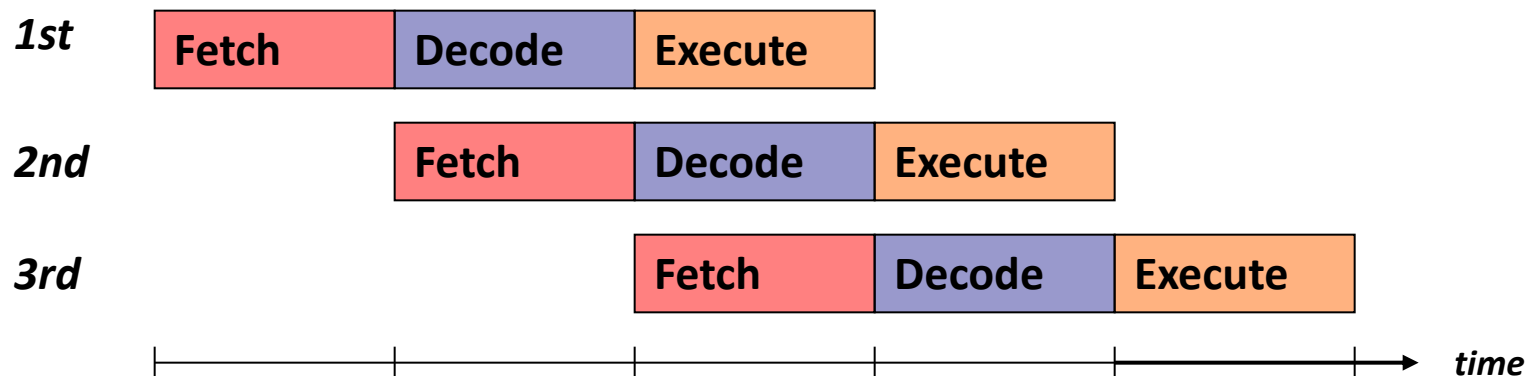


Instruction Pipeline

- The pipeline allows concurrent execution of multiple different instructions
 - execution of different stages of multiple instructions at the same time
- During a normal operation
 - while one instruction is being executed
 - the next instruction is being decoded
 - and a third instruction is being fetched from memory
 - allows effective throughput to increase to one instruction per clock cycle

Cortex-M3 Pipeline

- The Cortex-M3 Uses the 3-stage pipeline for instruction executions
 - Fetch \Rightarrow Decode \Rightarrow Execute
 - Pipeline design allows effective throughput to increase to one instruction per clock cycle
 - Allows the next instruction to be fetched while still decoding or executing the previous instructions

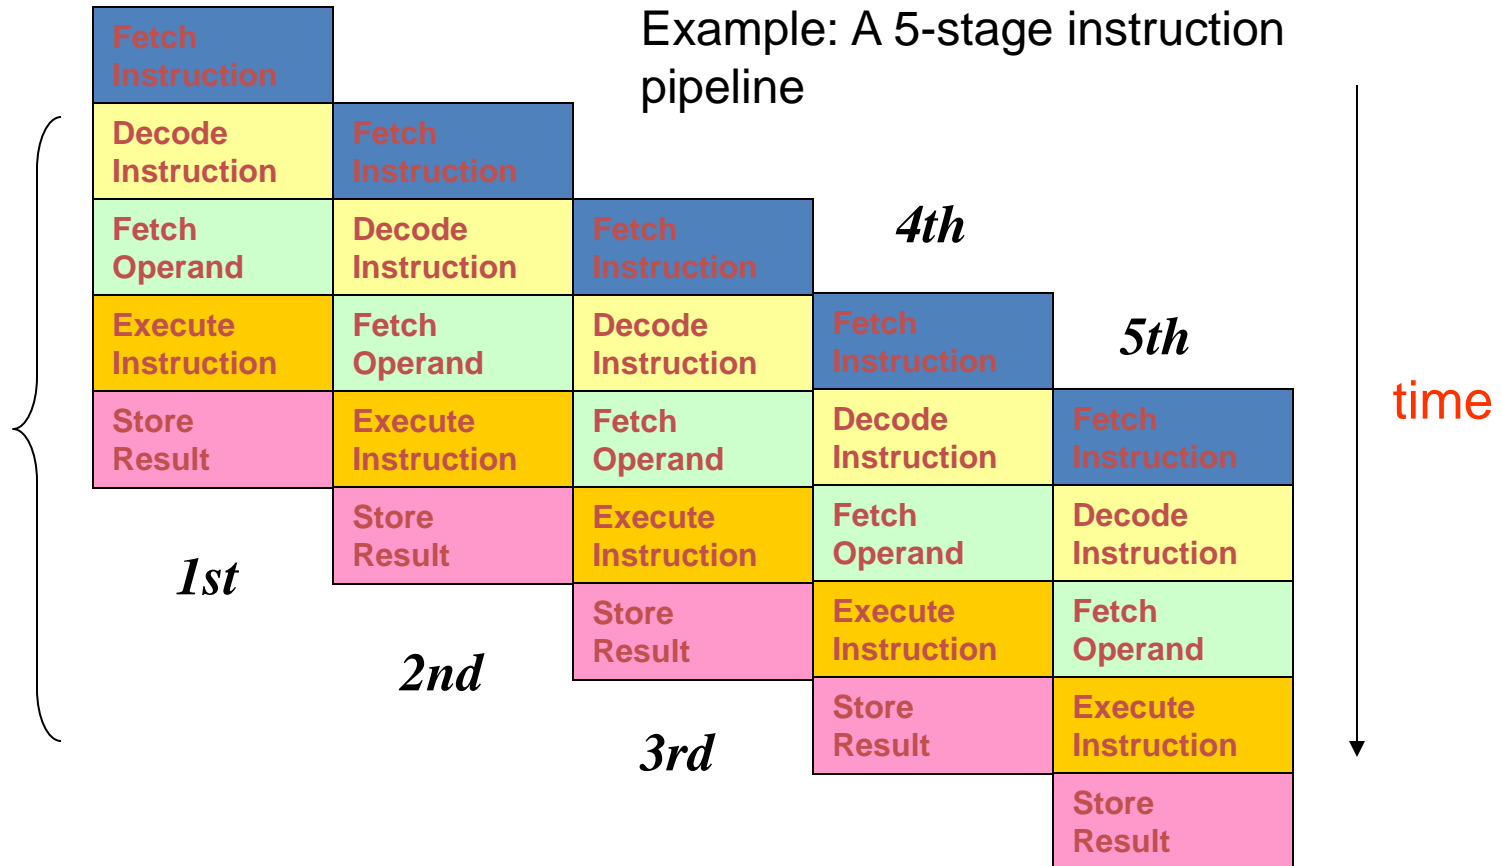


Pipelined Architecture

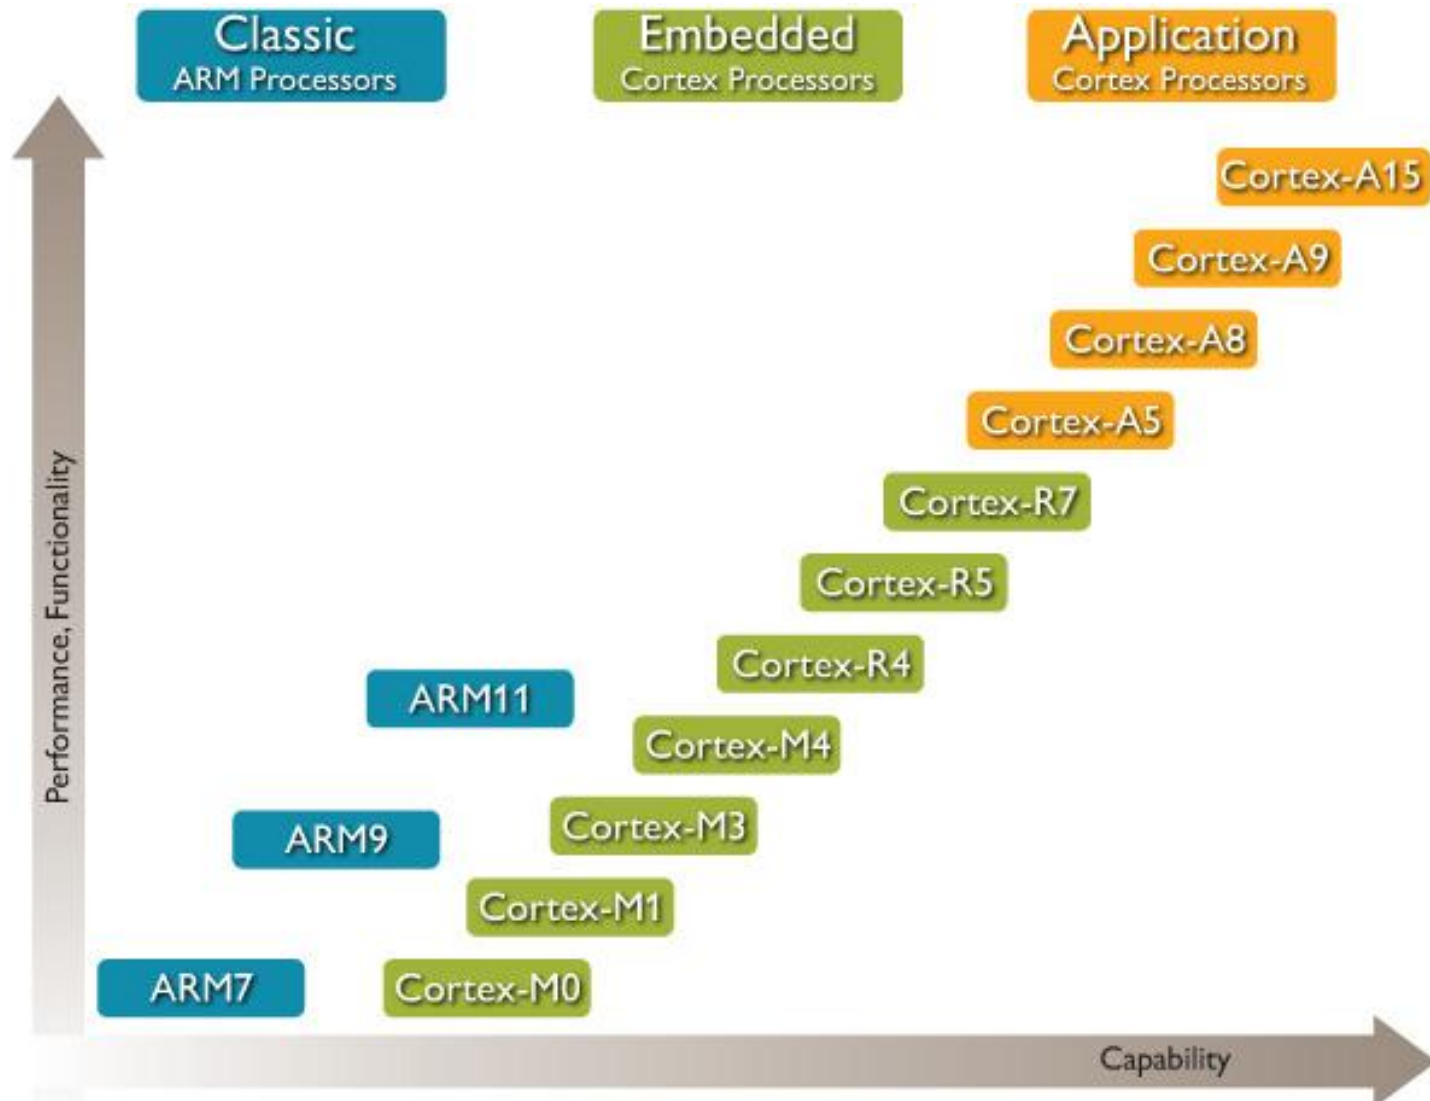
- A longer pipeline can also be used to further break down the operation carried out in the individual stage
 - simpler logic for each stage to increase system clock

Example: A 5-stage instruction pipeline

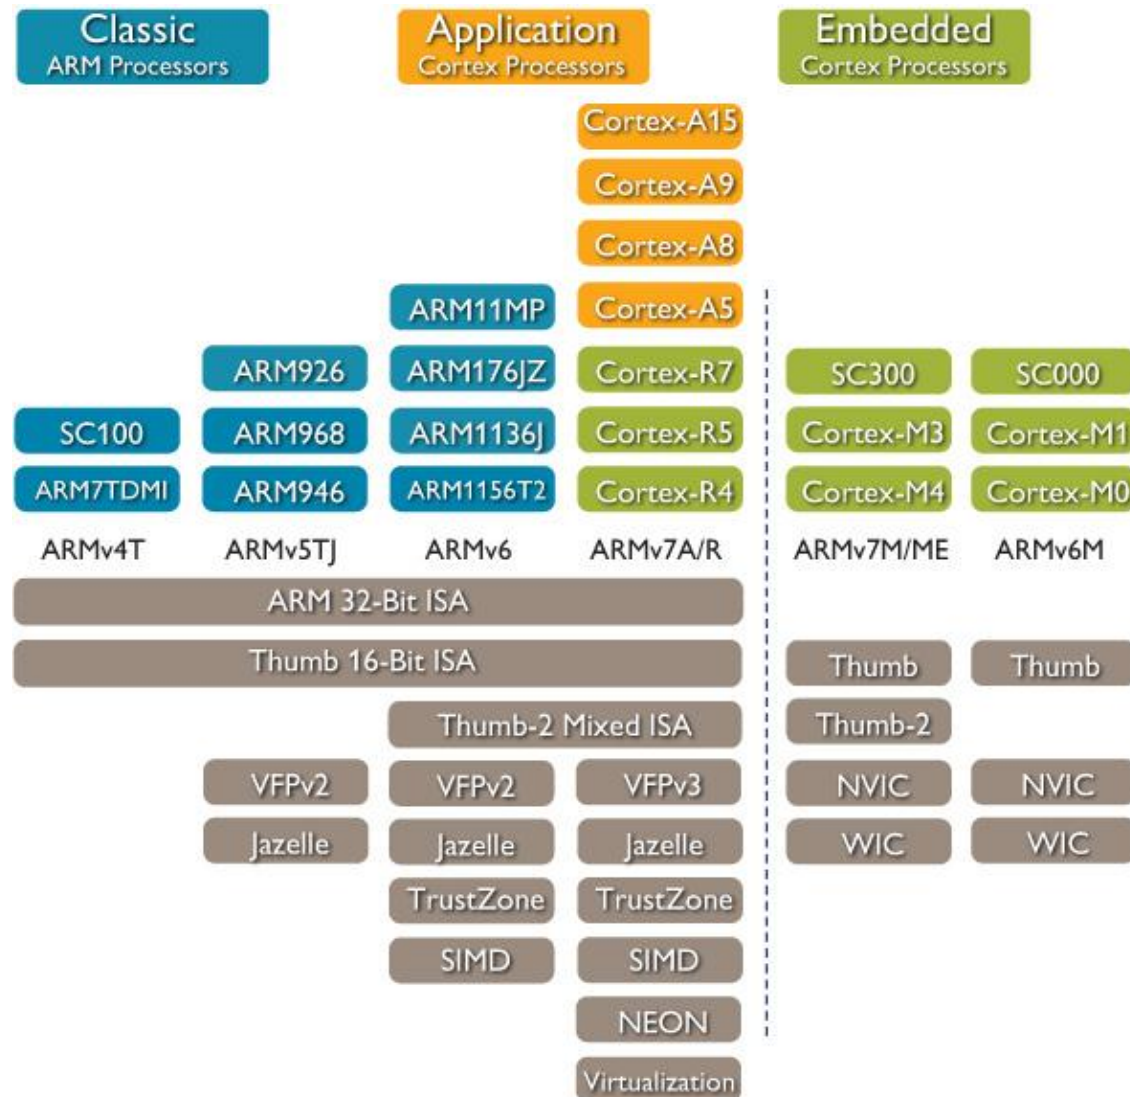
Parallel execution of multiple instructions



ARM Processors Families



ARM Processors Architectures



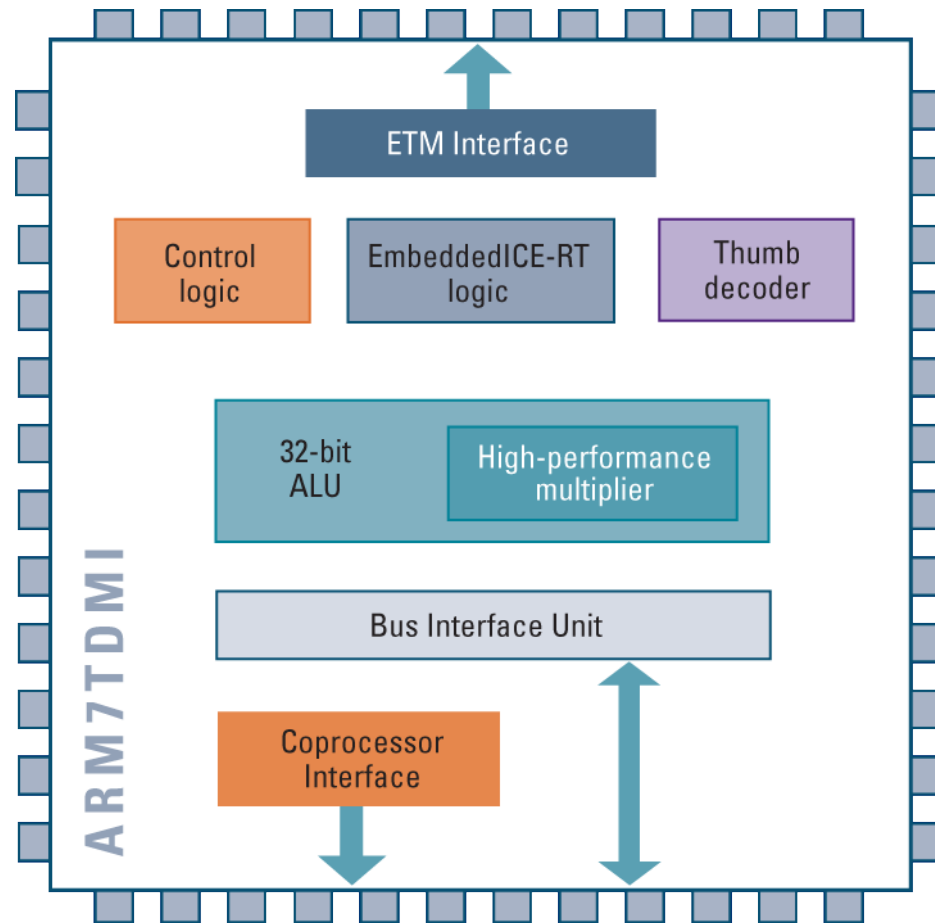
ARM Processors Architectures (2)

- Key attributes: Implementation size, performance, and very low power.
- Architectures types:
 - **ARMv4T** architecture introduced the 16-bit Thumb® instruction set alongside the 32-bit ARM instruction set.
 - **ARMv5TEJ** architecture introduced arithmetic support for digital signal processing (DSP) algorithms.
 - **ARMv6** architecture introduced an array of new features including the Single Instruction Multiple Data (SIMD) operations.
 - **ARMv7** architecture implements Thumb-2 technology.
 - **Cortex-A** implements a virtual memory system architecture based on an MMU, an optional NEON processing unit for multimedia applications and advanced hardware Floating Point.
 - **Cortex-R** – implements a protected memory system architecture based on an MPU (memory protection unit).
 - **Cortex-M** – Microcontroller profile designed for fast interrupt processing.

Classic ARM Processors

40

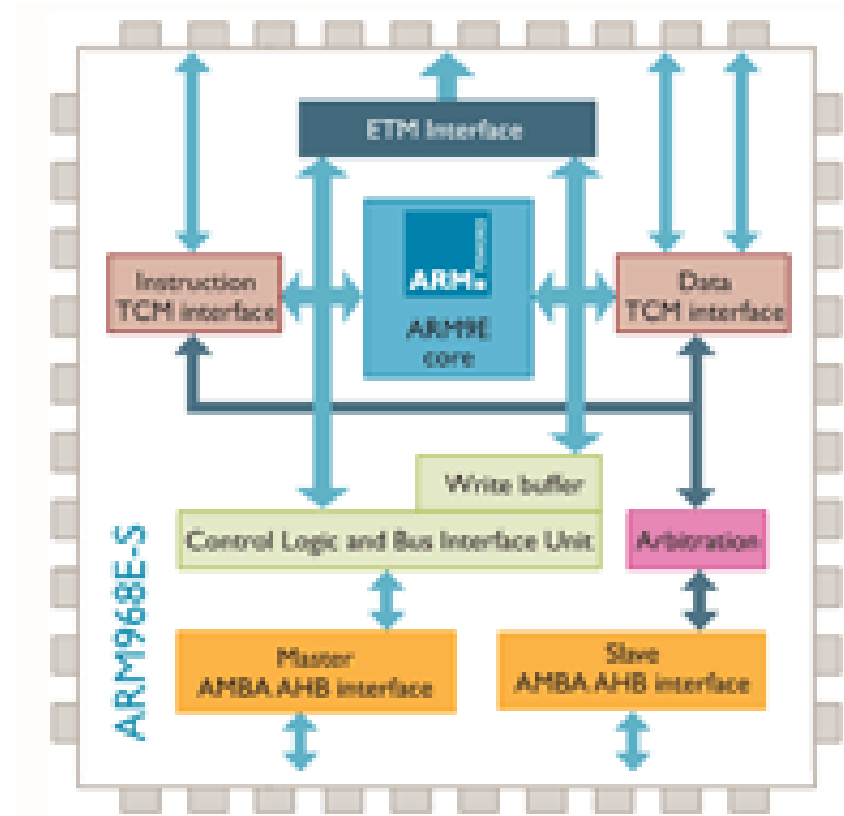
- ARM 7
 - Introduced in 1994.
 - More than 10 billion ARM7 processor family-based devices have powered a wide variety of applications.
 - Today is used for simple 32-bit devices.



Classic ARM Processors (2)

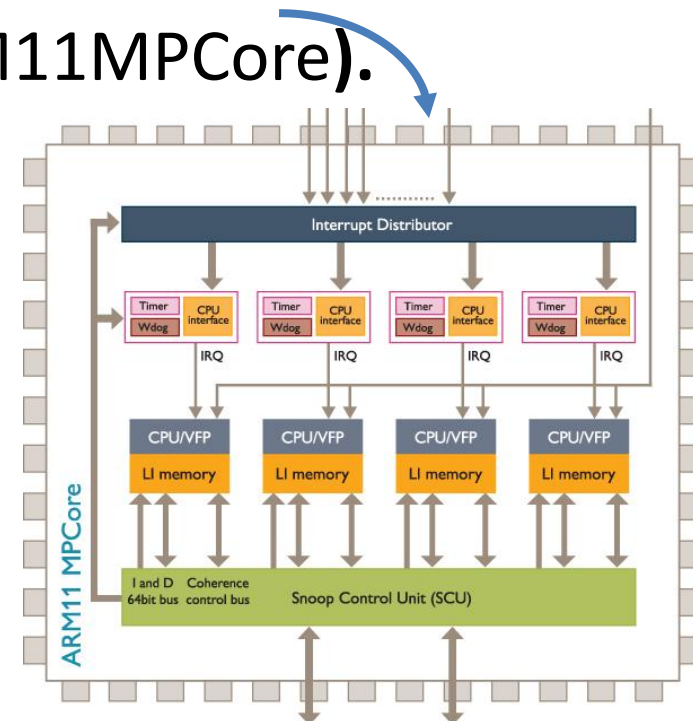
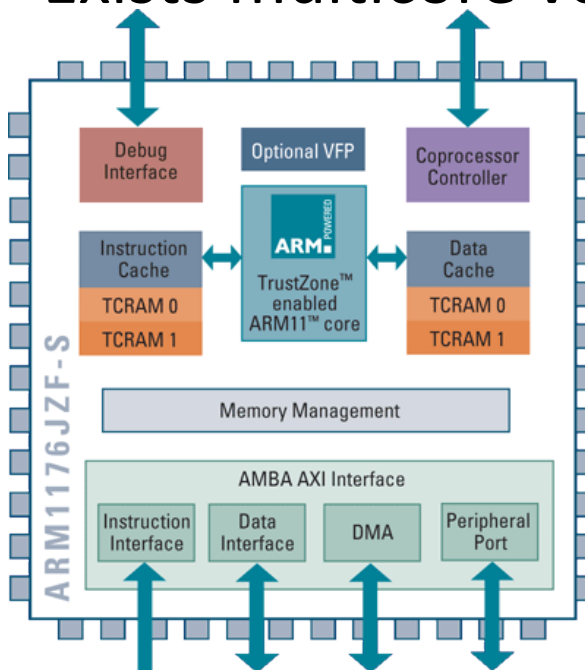
41

- ARM 9
 - Is the most popular ARM processor family ever.
 - Over 5 Billion ARM9 processors have been shipped so far.
 - Successfully deployed across a wide range of applications.



Classic ARM Processors (3)

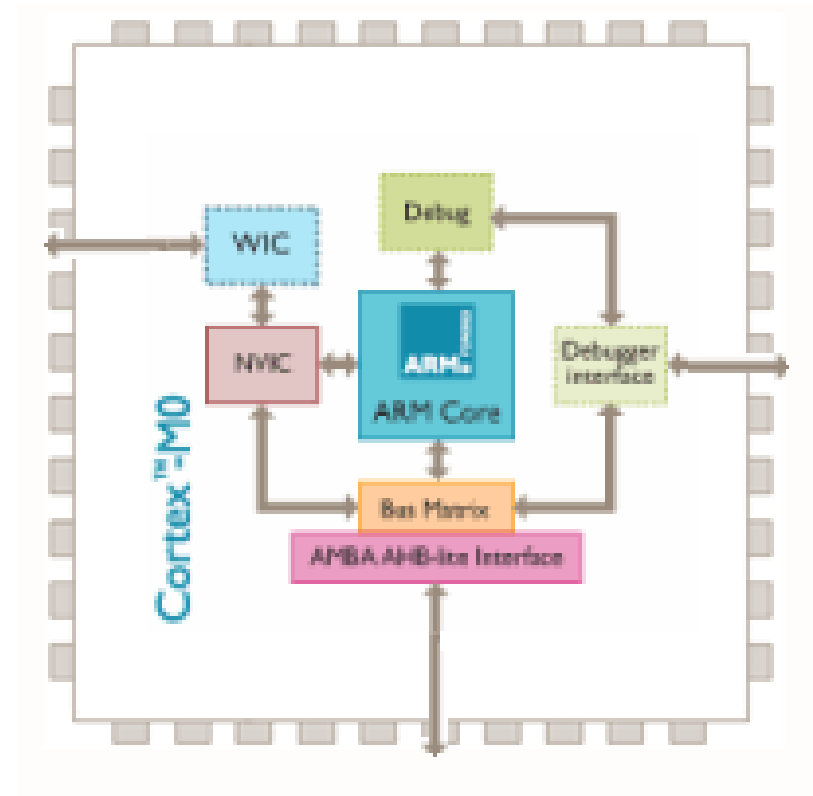
- ARM 11
 - Powers many today's smartphones.
 - Extreme low power.
 - Exists multicore version (ARM11MPCore).



Embedded ARM Cortex Processors

43


- Cortex M0:
 - Ultra low gate count (less than 12 K gates).
 - Ultra low-power (3 μ W/MHz).
 - 32-bit processor.



Embedded ARM Cortex Processors (2)

44

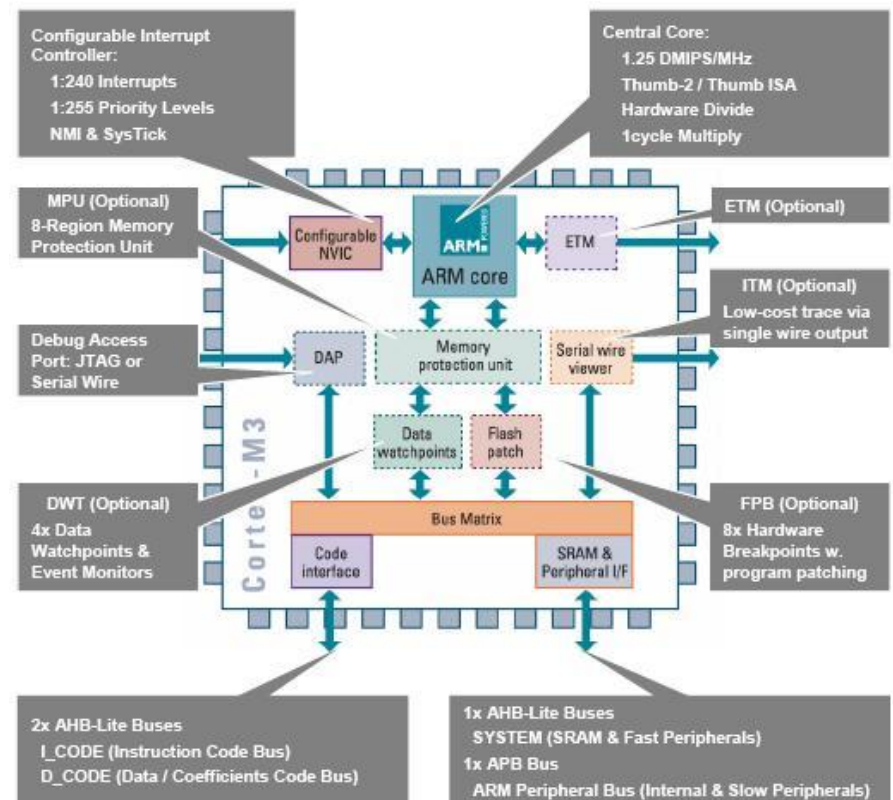
- Cortex M1:
 - The first ARM processor designed specifically for implementation in FPGAs.
 - Supports all major FPGA vendors.
 - Easy migration path from FPGA to ASIC.



Embedded ARM Cortex Processors (3)

45

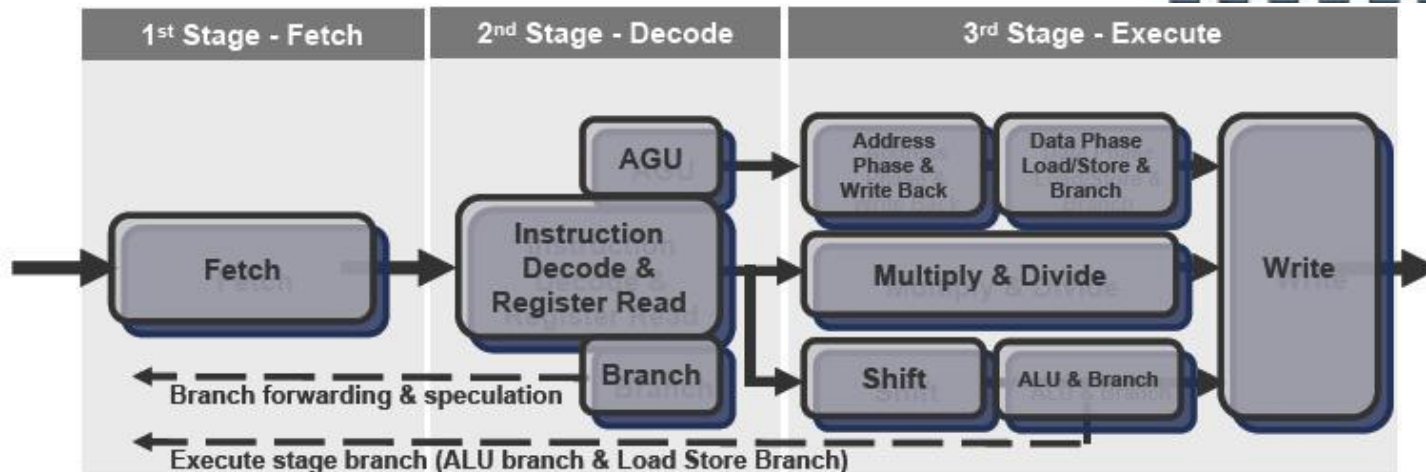
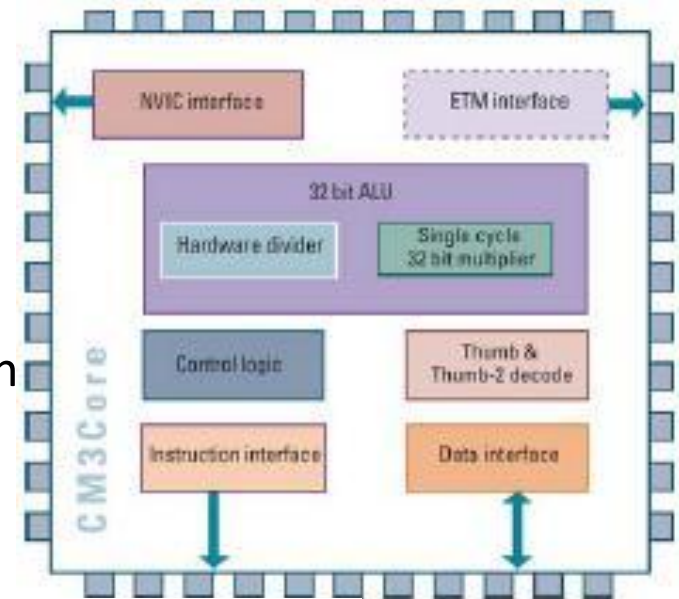
- Cortex M3:
 - The mainstream ARM processor for microcontroller applications.
 - High performance and energy efficiency.



Cortex M3 Central Core

46

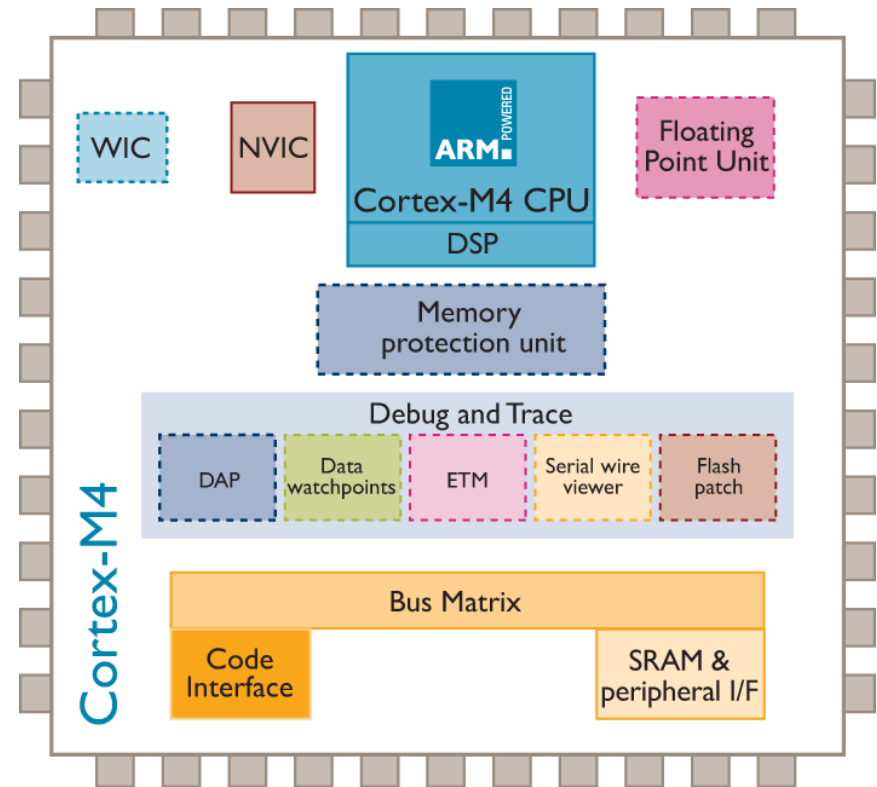
- Harvard architecture:
 - Separate Instruction & Data buses enable parallel fetch & store.
- Advanced 3-Stage Pipeline:
 - Includes Branch Forwarding & Speculation
- Additional Write-Back via Bus Matrix.



Embedded ARM Cortex Processors (4)

47

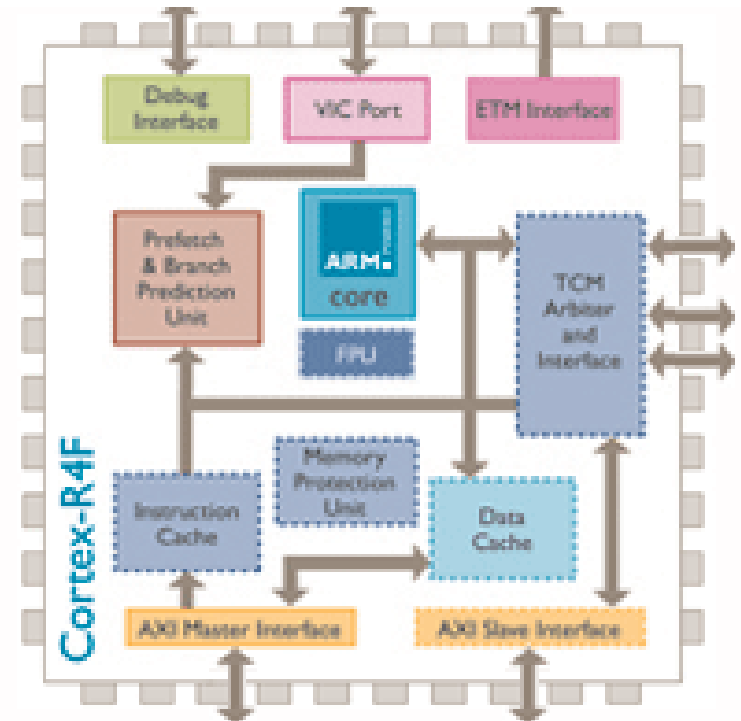
- Cortex M4:
 - The latest embedded processor for DSP.



Embedded ARM Cortex Processors (5)

48

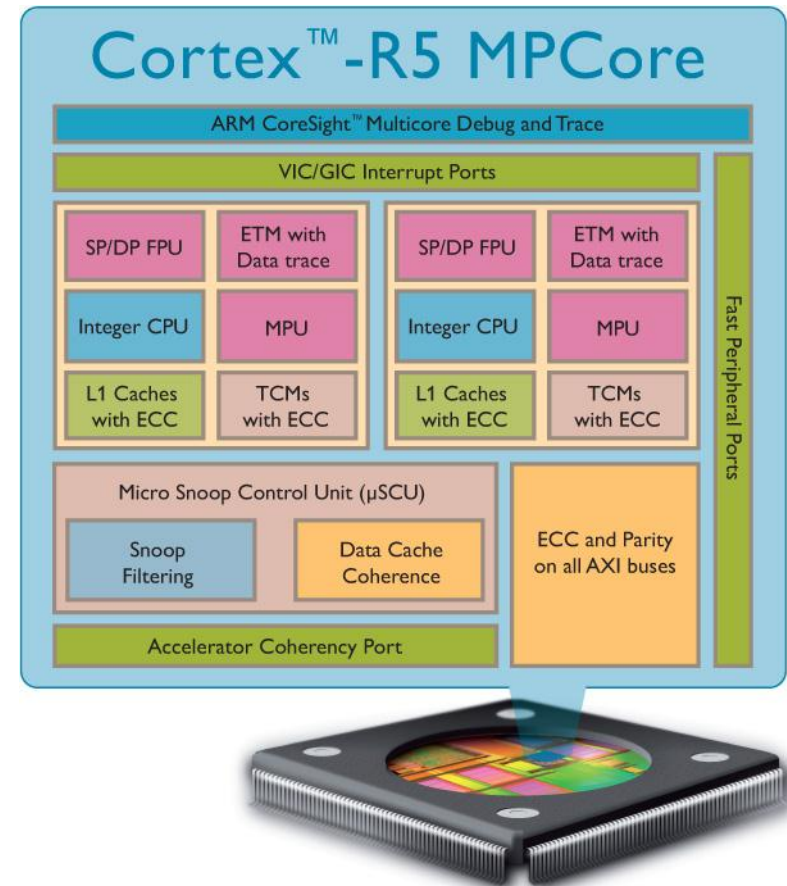
- Cortex R4:
 - First embedded real-time processor based on the ARMv7-R architecture.
 - For high-volume deeply-embedded System-on-Chip applications:
 - Hard disk.
 - Drive controllers.
 - Wireless baseband processors.
 - Electronic control units for automotive systems.



Embedded ARM Cortex Processors (6)

49

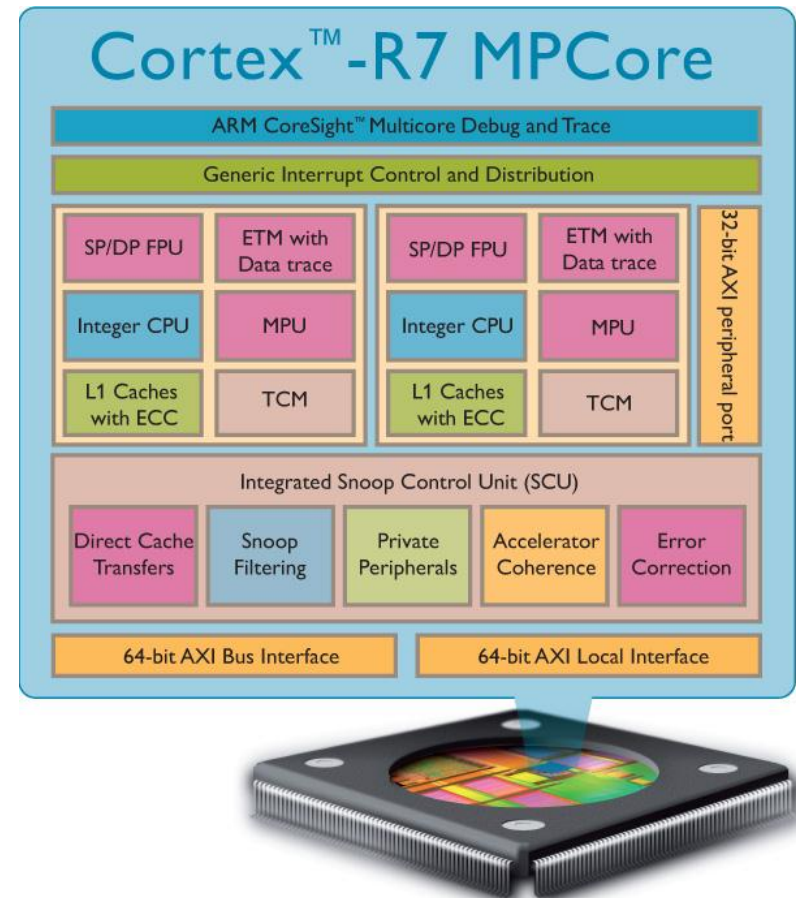
- Cortex R5:
 - Extends the feature set of the Cortex-R4.
 - Enables:
 - Higher levels of system performance.
 - Increased efficiency and reliability.
 - Enhanced error management in dependable real-time systems.



Embedded ARM Cortex Processors (7)

50

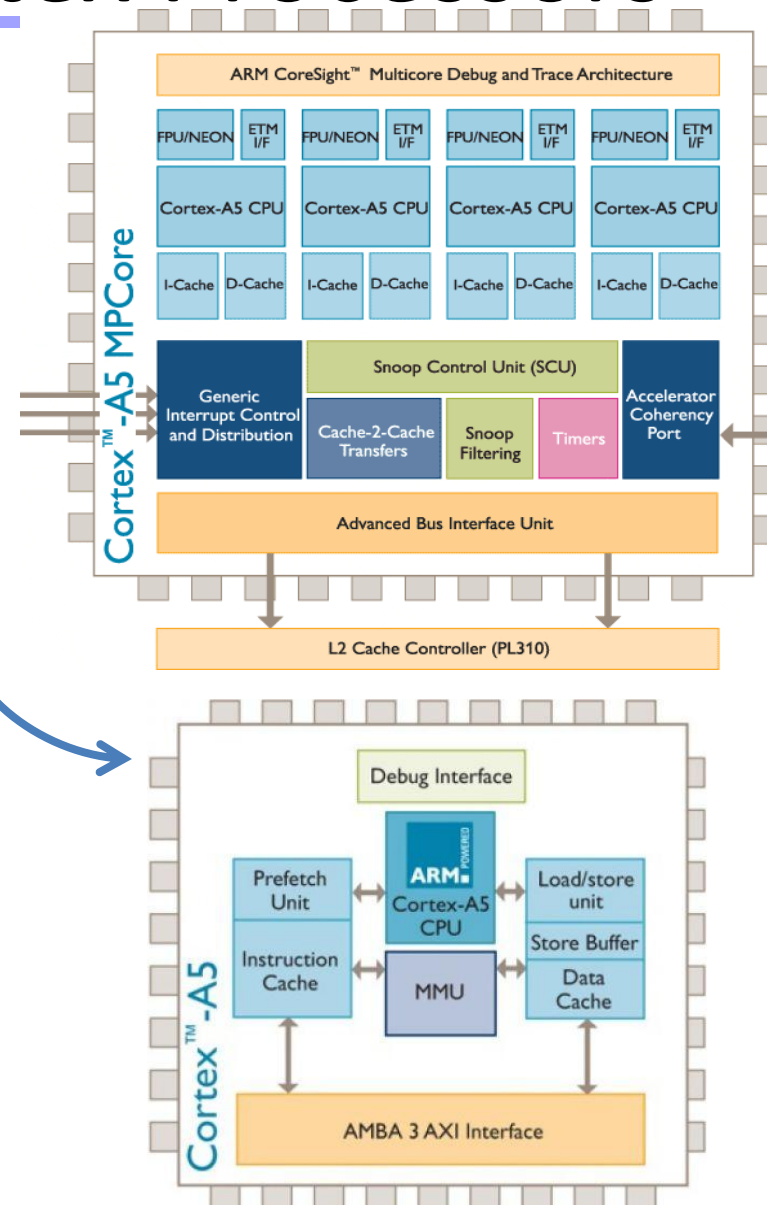
- Cortex R7:
 - High-performance dual-core.
 - It is the highest performing Cortex-R series processor.



Application ARM Cortex Processors

51

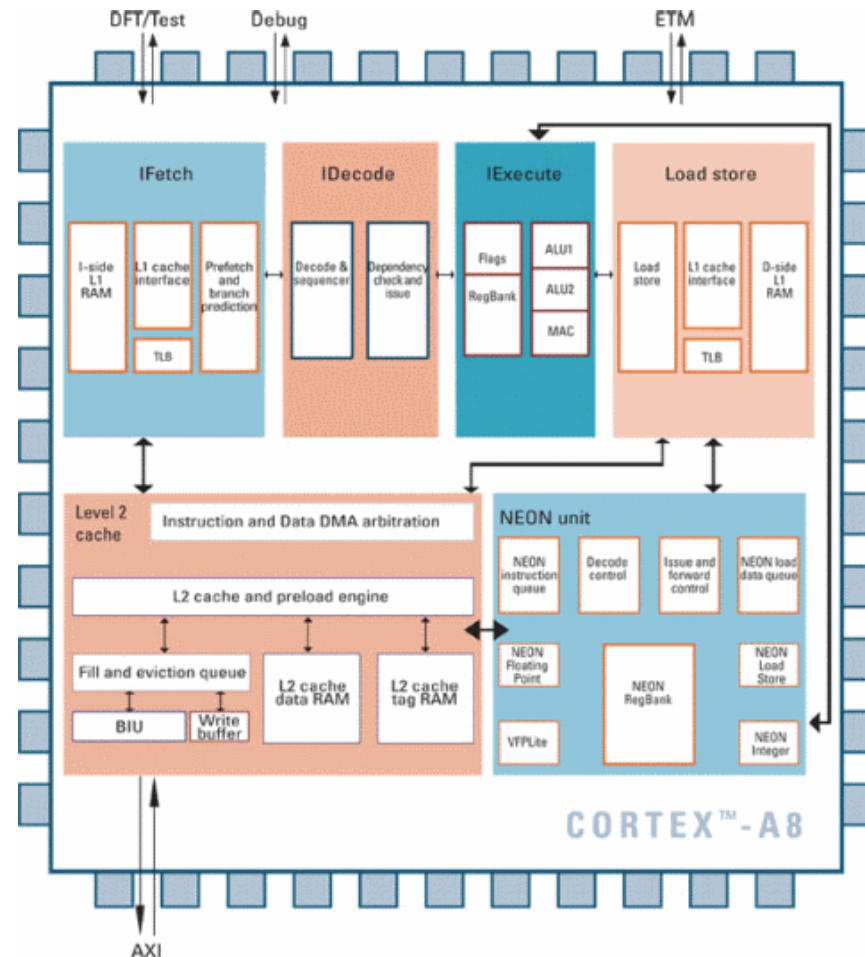
- Cortex A5:
 - Delivers high end features to power and cost sensitive applications.
 - Single core version also available.
 - Suitable for:
 - From entry level smartphones, low cost handsets and smart mobile devices.
 - To pervasive embedded, consumer and industrial devices.



Application ARM Cortex Processors (2)

52

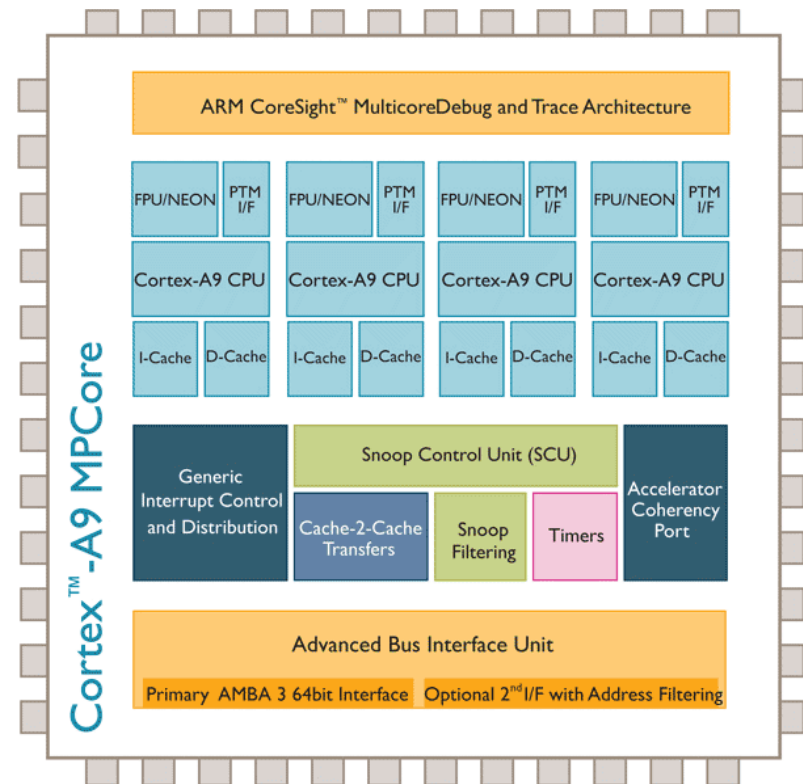
- Cortex A8:
 - Single core only.
 - Can scale in speed from 600MHz to greater than 1GHz.
 - Suitable for high-end feature phones, netbooks, DTVs, printers and automotive-infotainment.



Application ARM Cortex Processors (3)

53

- Cortex A9:
 - From 1 to 4 cores.
 - Extremely high levels of performance and power efficiency.
 - Ideal solution for designs requiring high performance in low power or thermally constrained cost-sensitive devices.



Application ARM Cortex Processors (4)

54

- Cortex A15:
 - Ultra low-power.
 - Suitable for:
 - Advanced Smartphones.
 - Mobile Computing.
 - High-end Digital Home Entertainment.
 - Wireless Infrastructure.
 - Low-power Servers.

